

CHERRY: CHECKPOINTED EARLY RESOURCE RECYCLING

José F. Martínez¹, Jose Renau²

Michael C. Huang³, Milos Prvulovic², and Josep Torrellas²

1



2



3





MOTIVATION

- **Problem:** Limited processor resources
- **Goal:** More efficient use by aggressive recycling
- **Opportunity:** Resources reserved until retirement
- **Solution:** Decouple recycling from retirement





CURRENT PROCESSORS

LDQ Entry

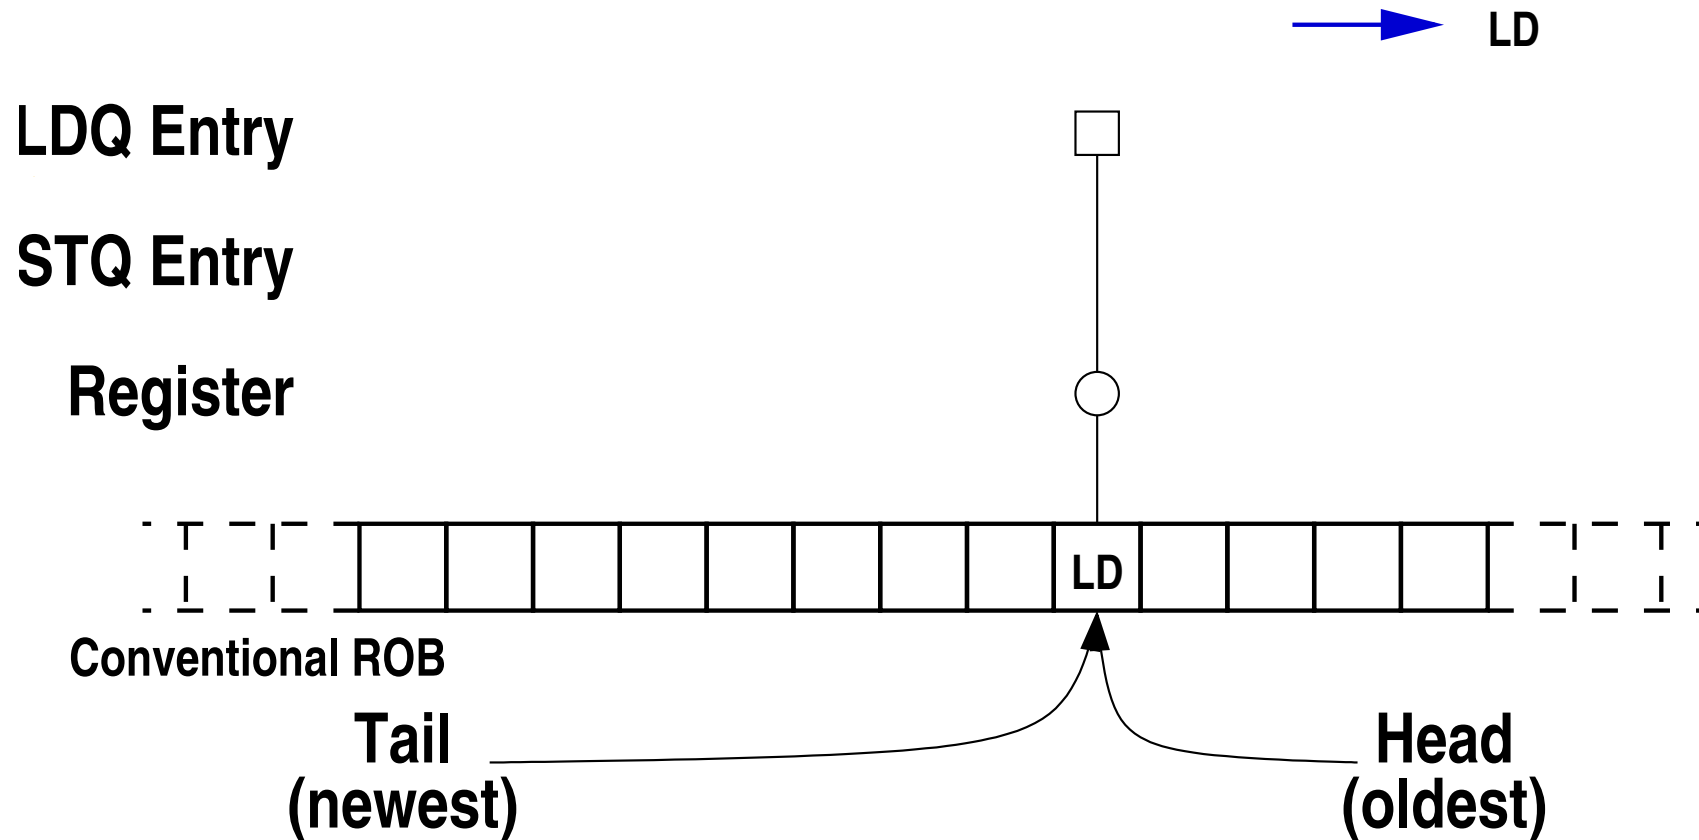
STQ Entry

Register



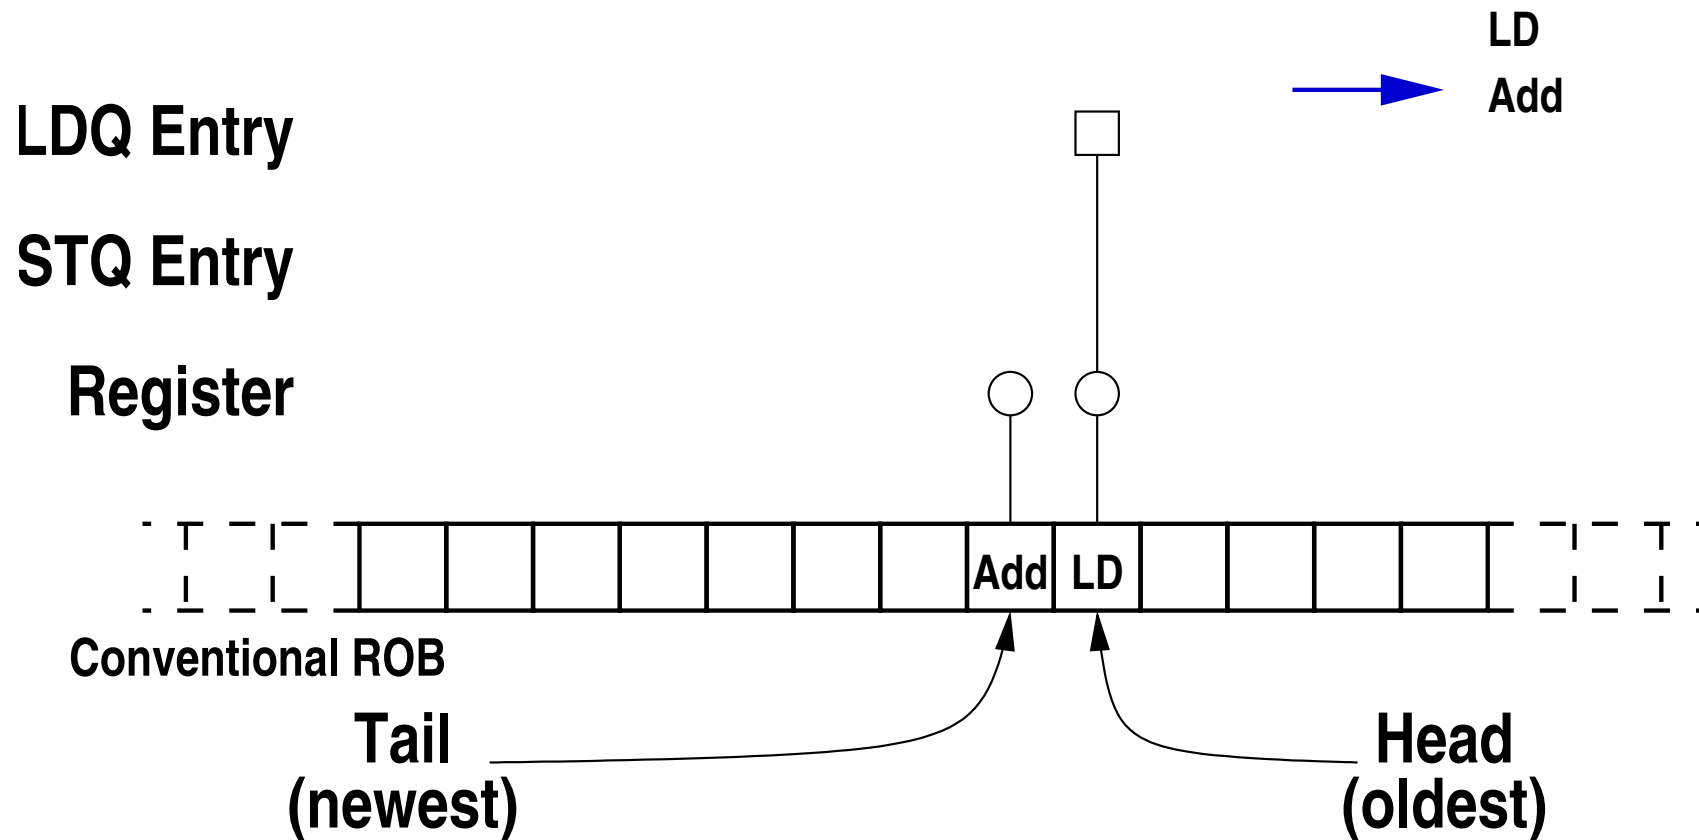



CURRENT PROCESSORS





CURRENT PROCESSORS






CURRENT PROCESSORS


LDQ Entry

STQ Entry

Register

LD
Add
BR






CURRENT PROCESSORS


LDQ Entry

STQ Entry

Register

LD
Add
BR
ST






CURRENT PROCESSORS


LDQ Entry

STQ Entry

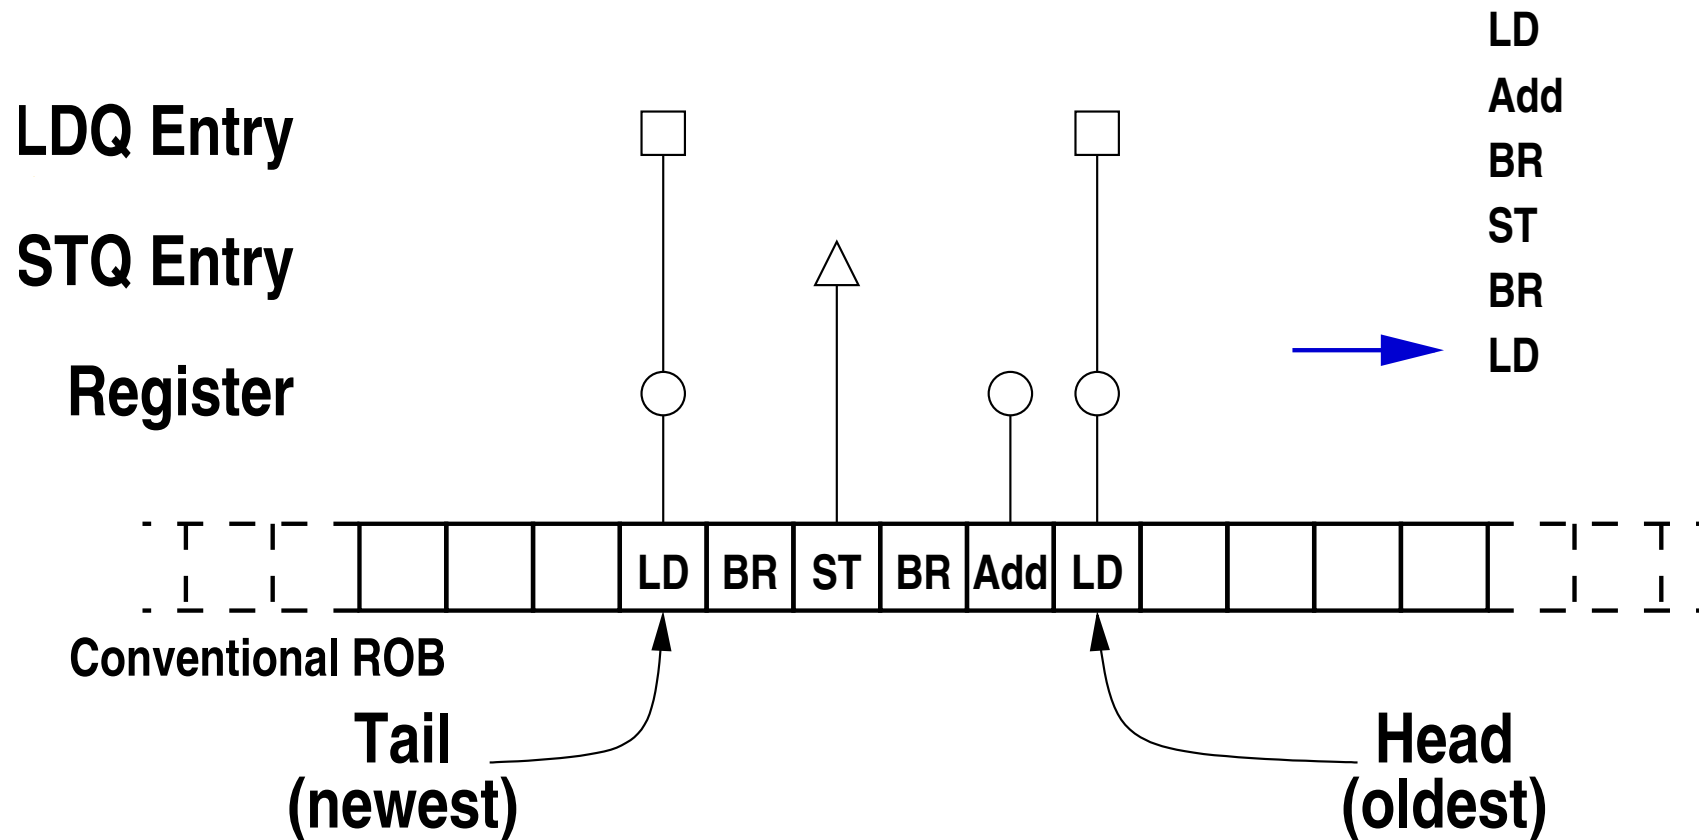
Register


LD
Add
BR
ST
BR



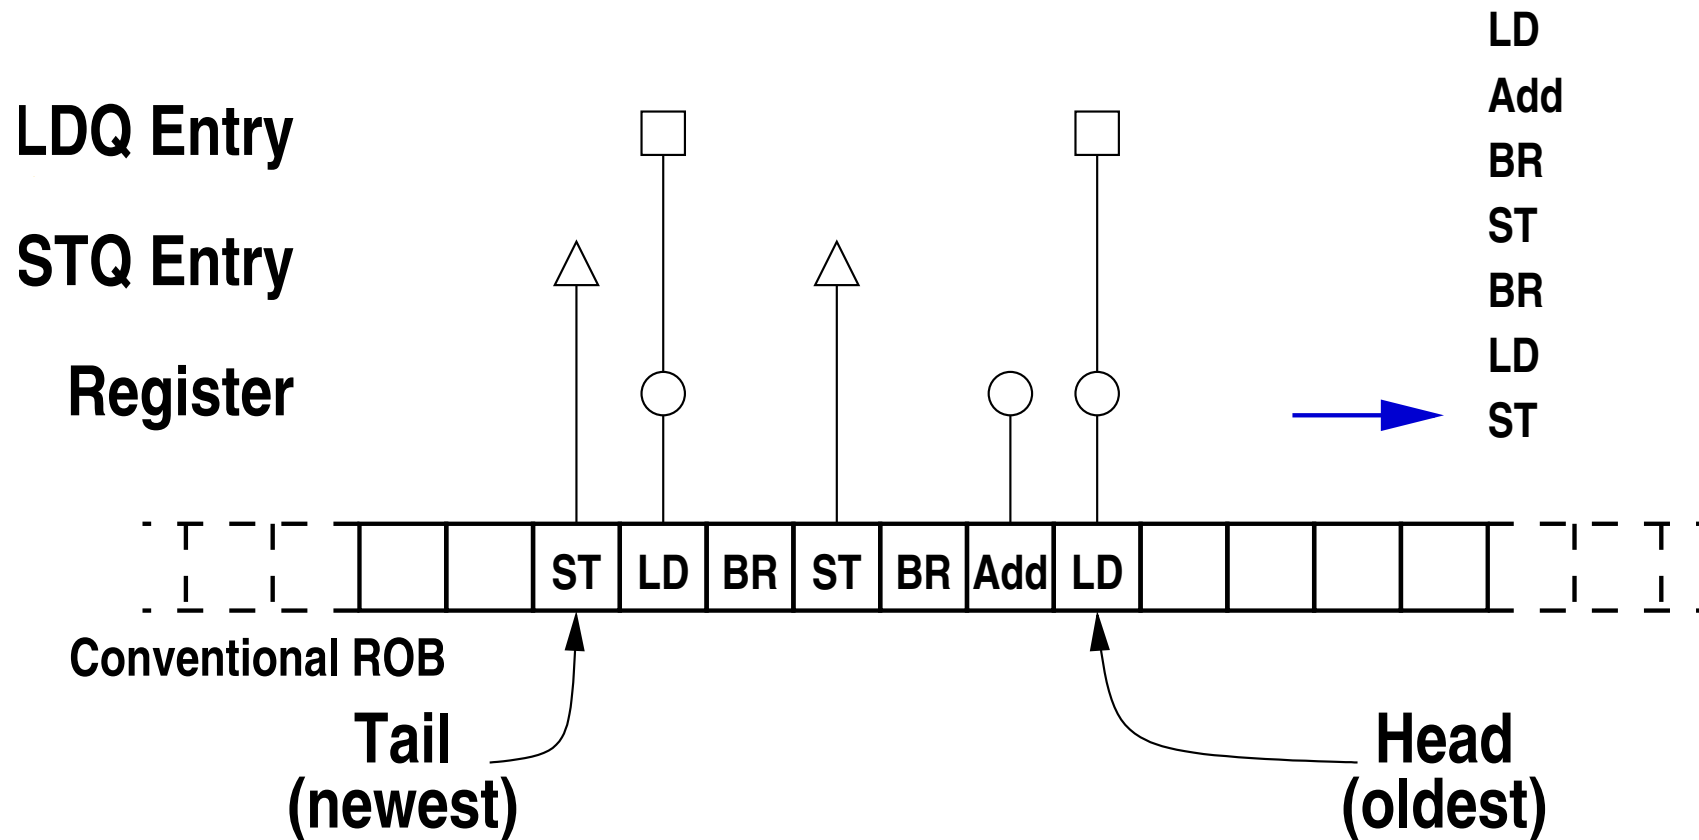



CURRENT PROCESSORS



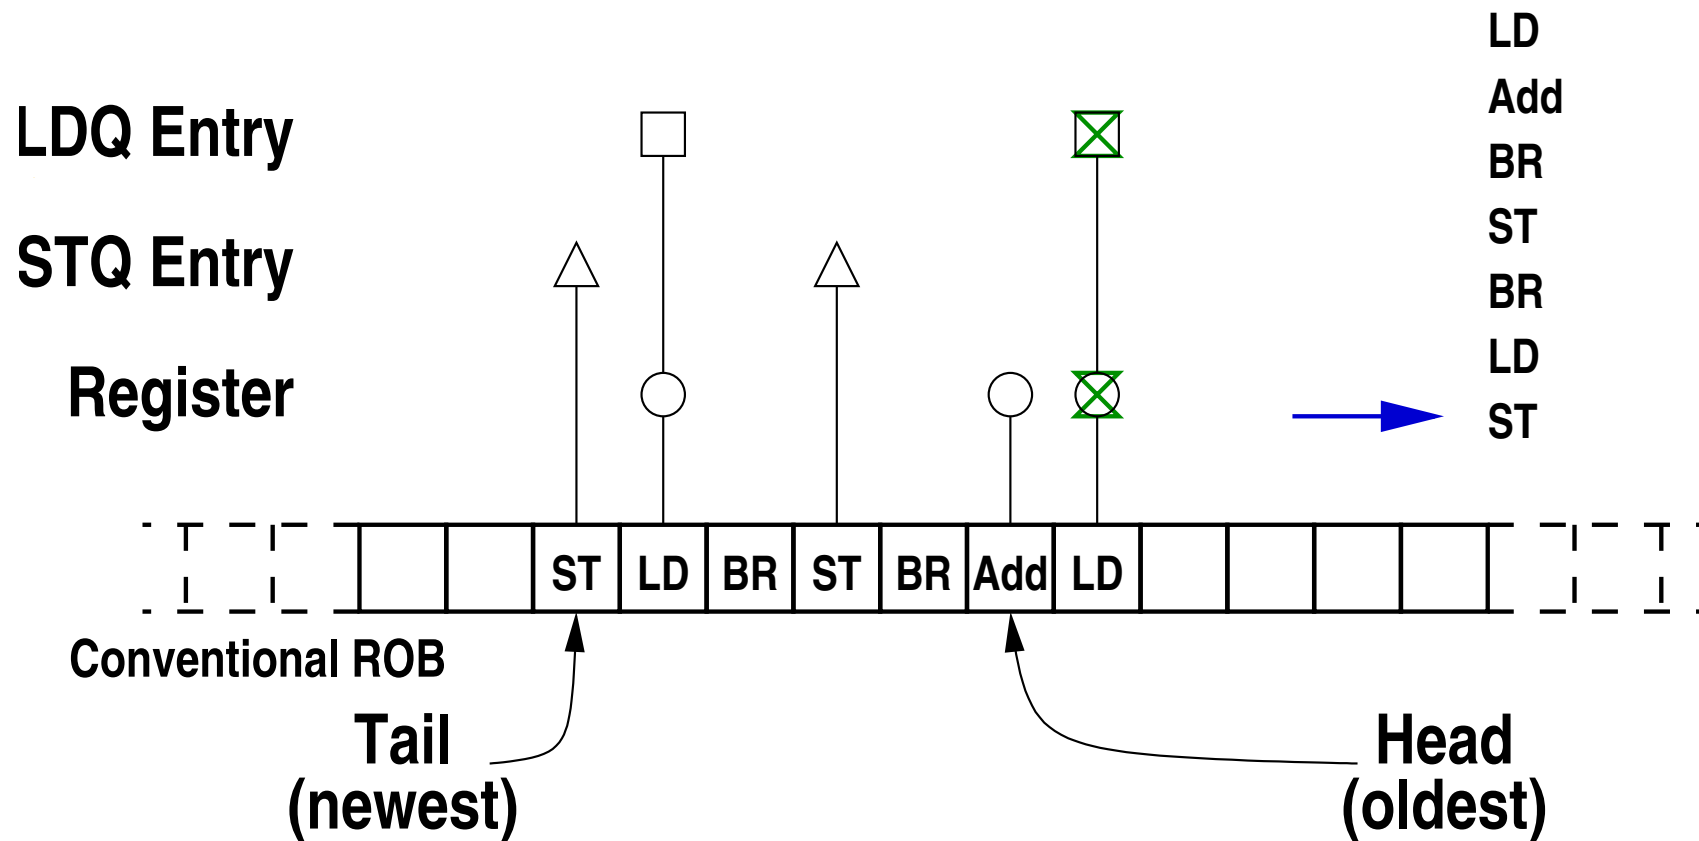



CURRENT PROCESSORS



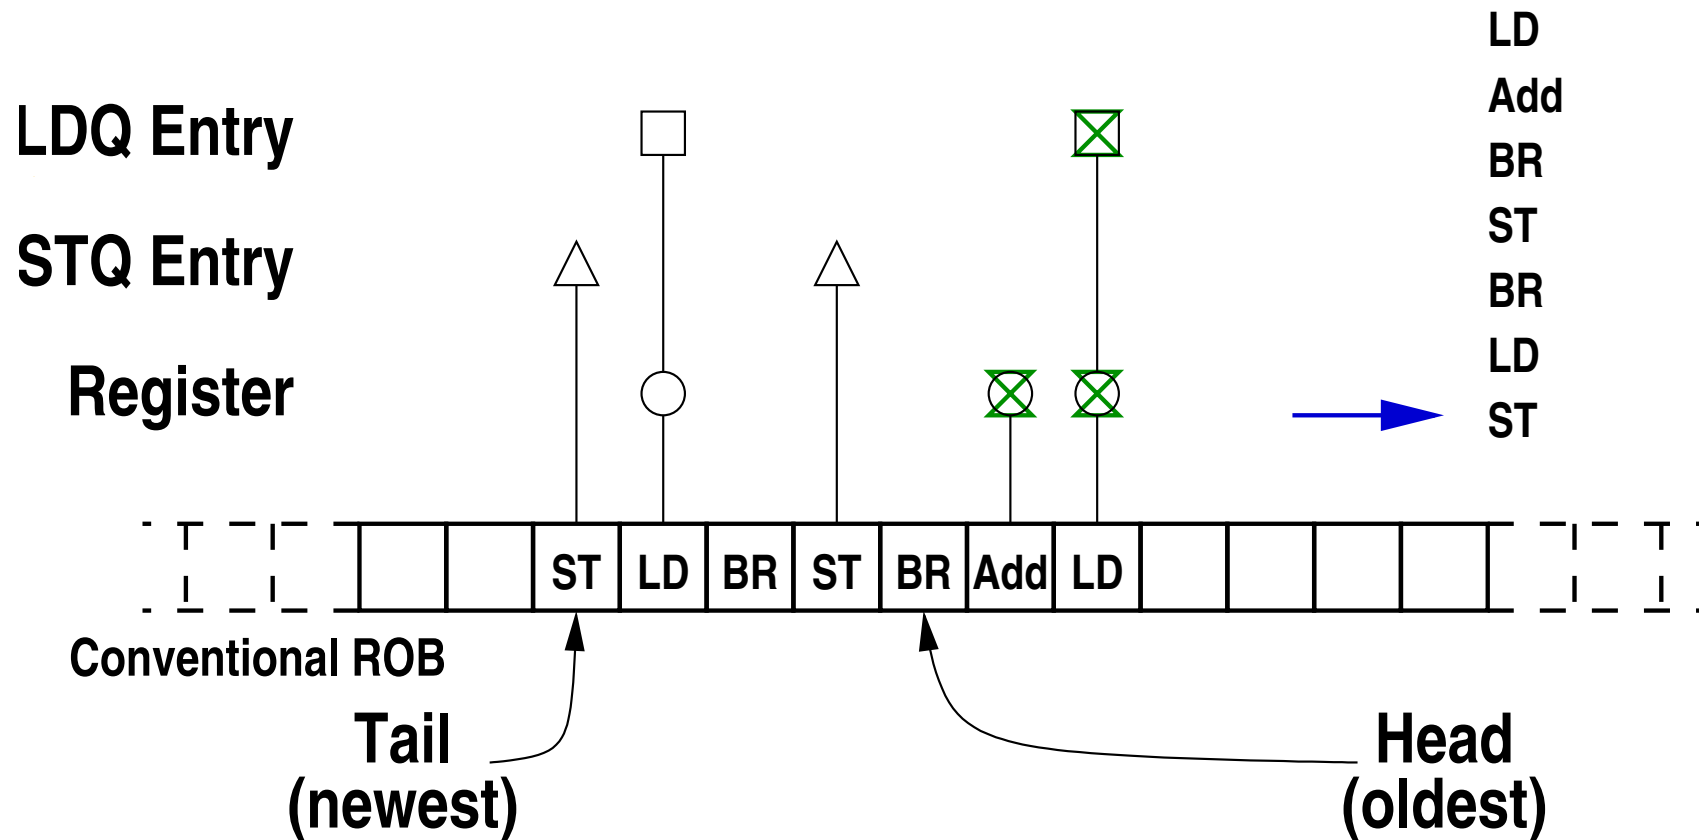


CURRENT PROCESSORS



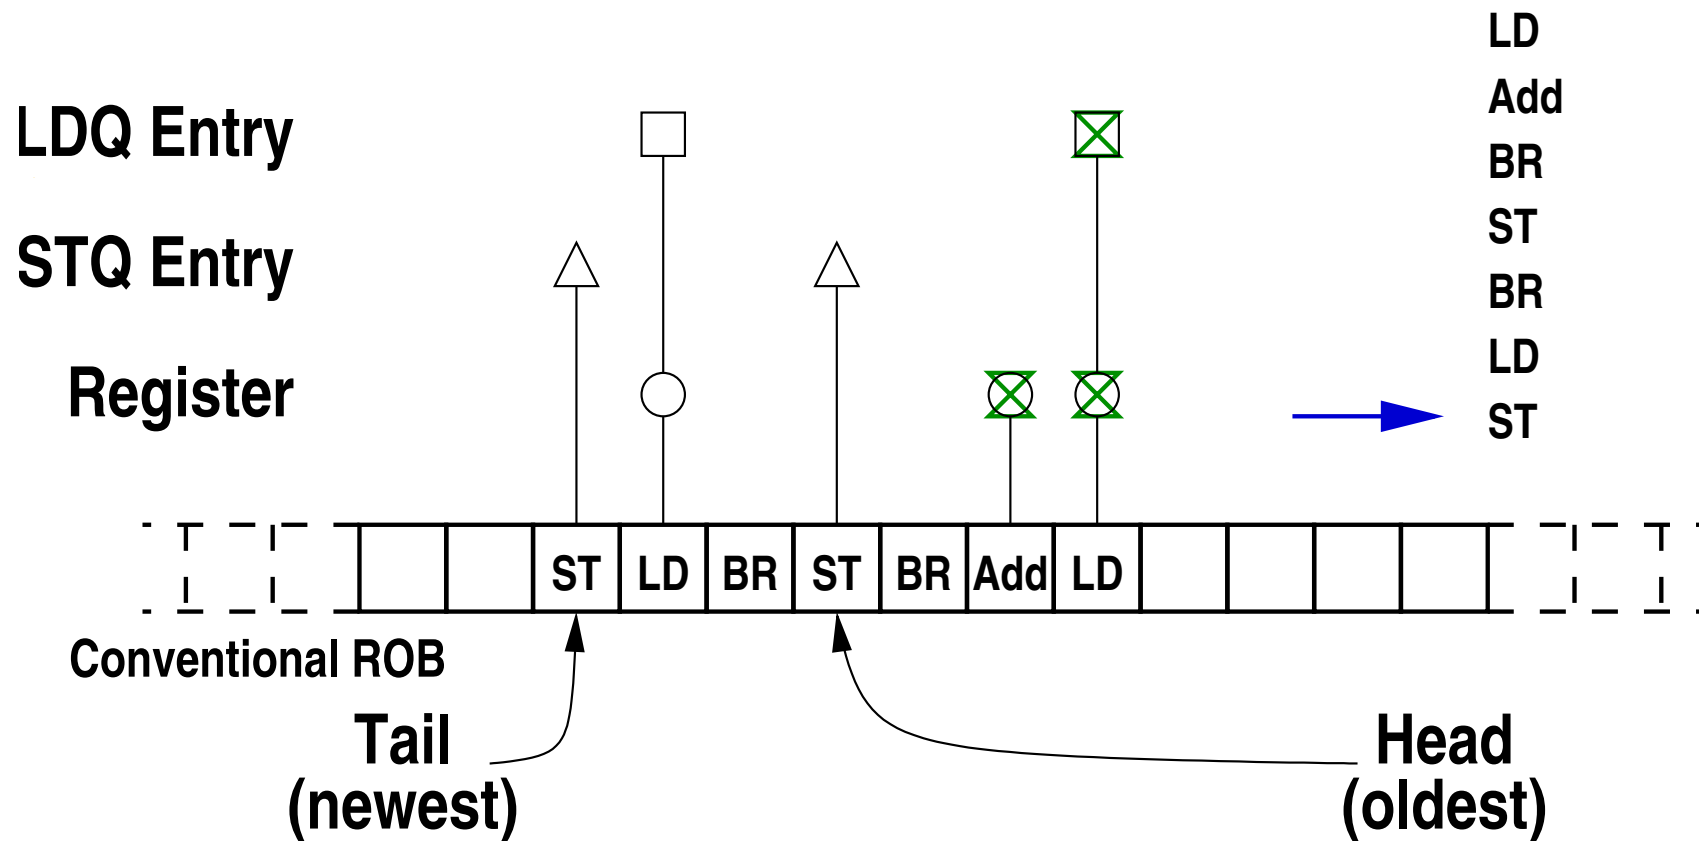



CURRENT PROCESSORS






CURRENT PROCESSORS







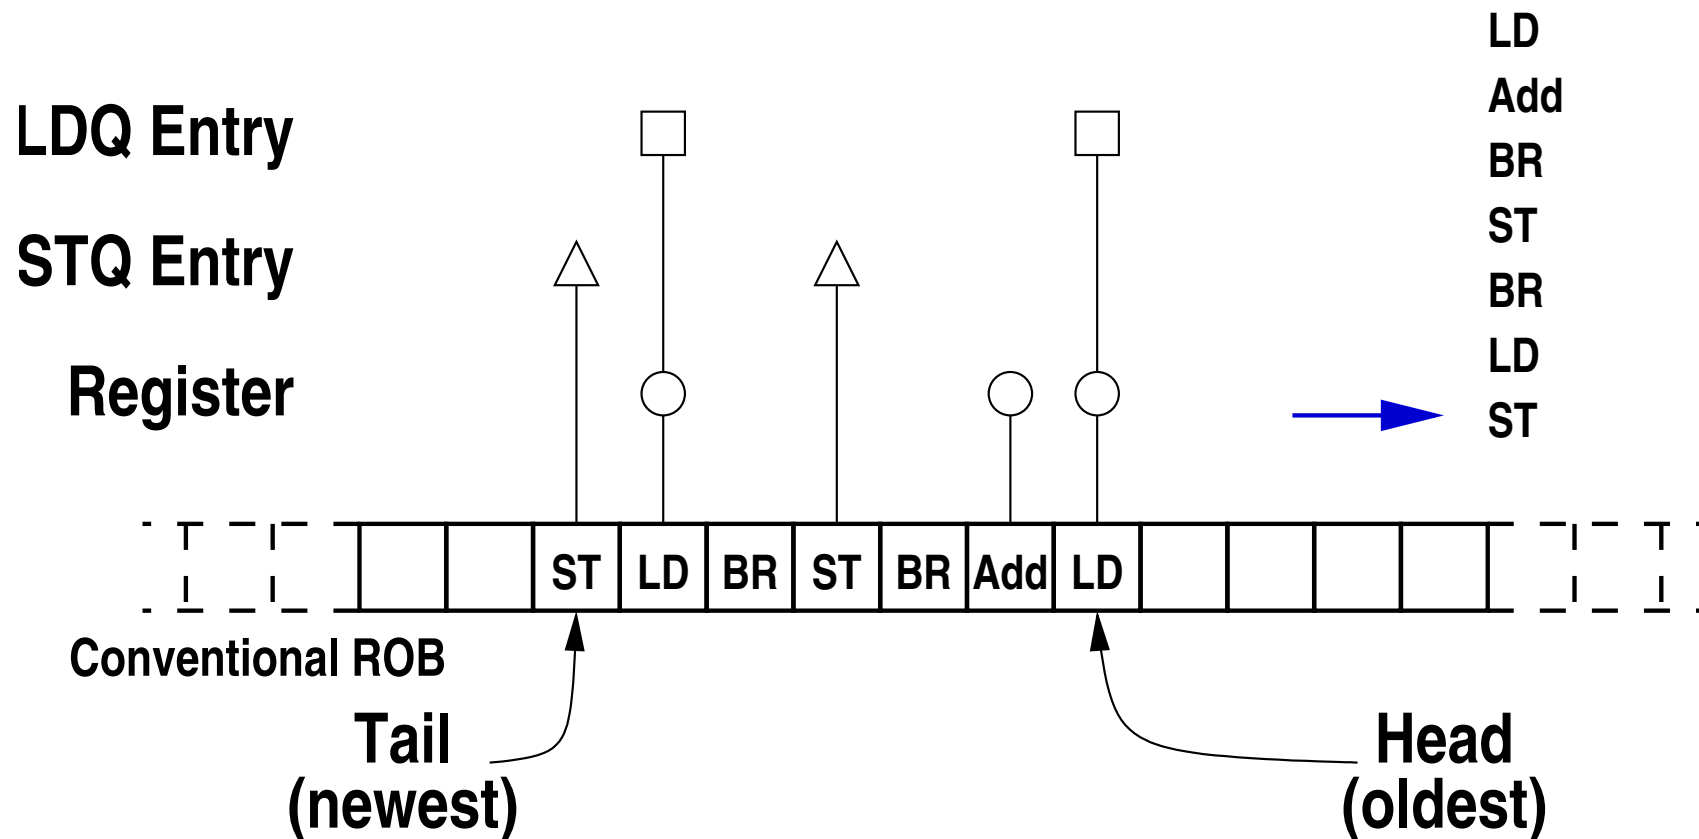
PROPOSAL: EARLY RECYCLING

- Decouple resource recycling from instruction retirement
 - Recycle when no longer needed
- Targeted resources:
 - Load queue entries
 - Store queue entries
 - Physical registers
- Potential when targeted resources are unlimited:
 - 1.12 speedup in SPECint2000
 - 1.32 speedup in SPECfp2000



EARLY RECYCLING

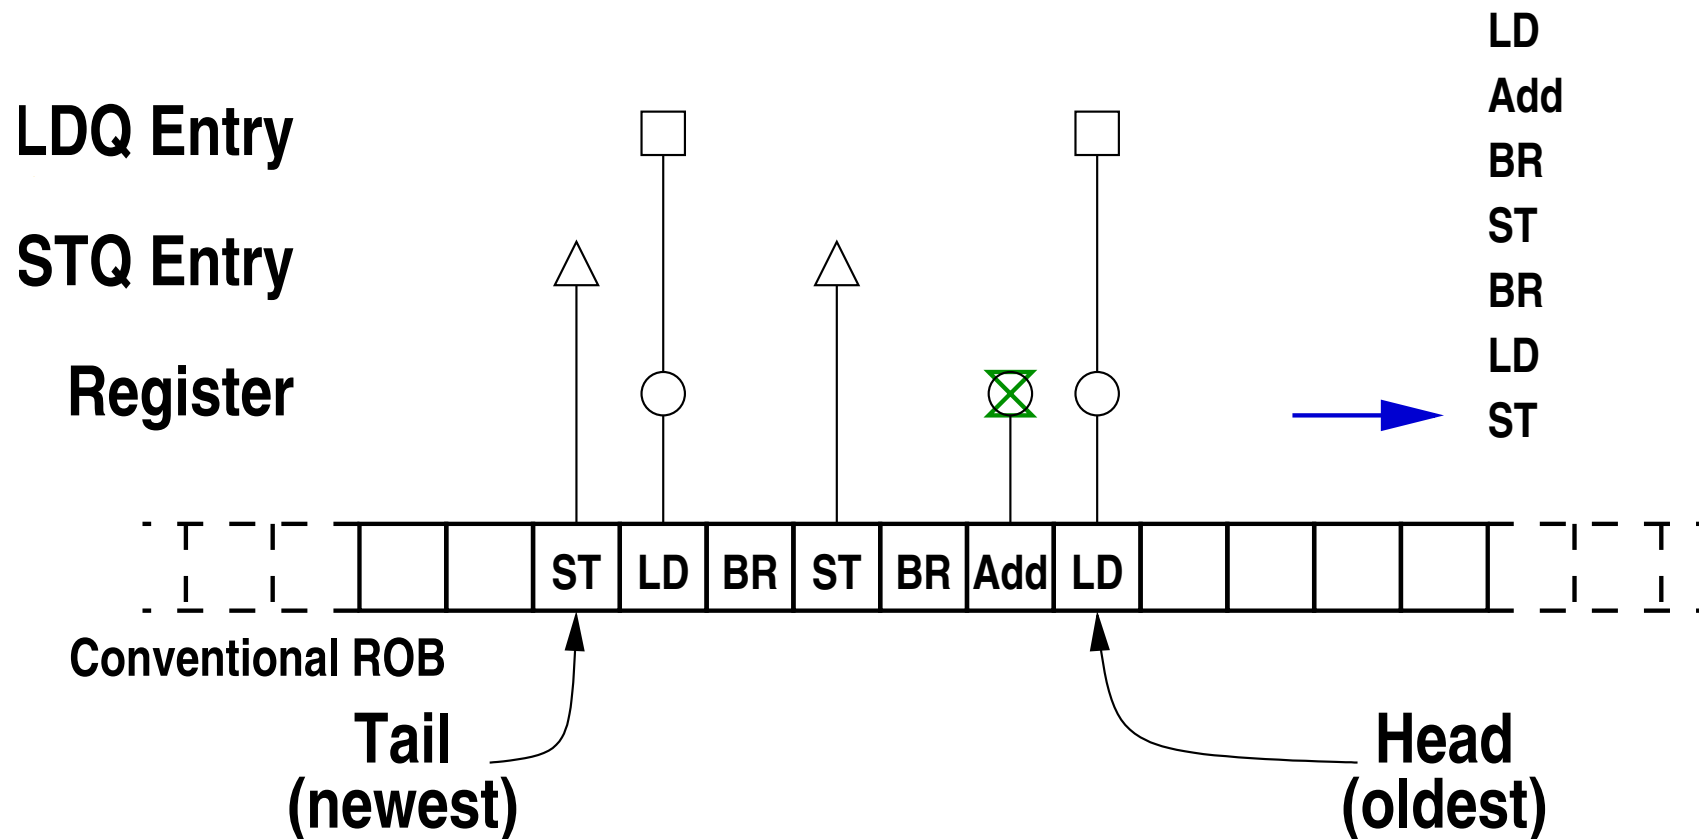
- Decouple resource recycling from retirement





EARLY RECYCLING

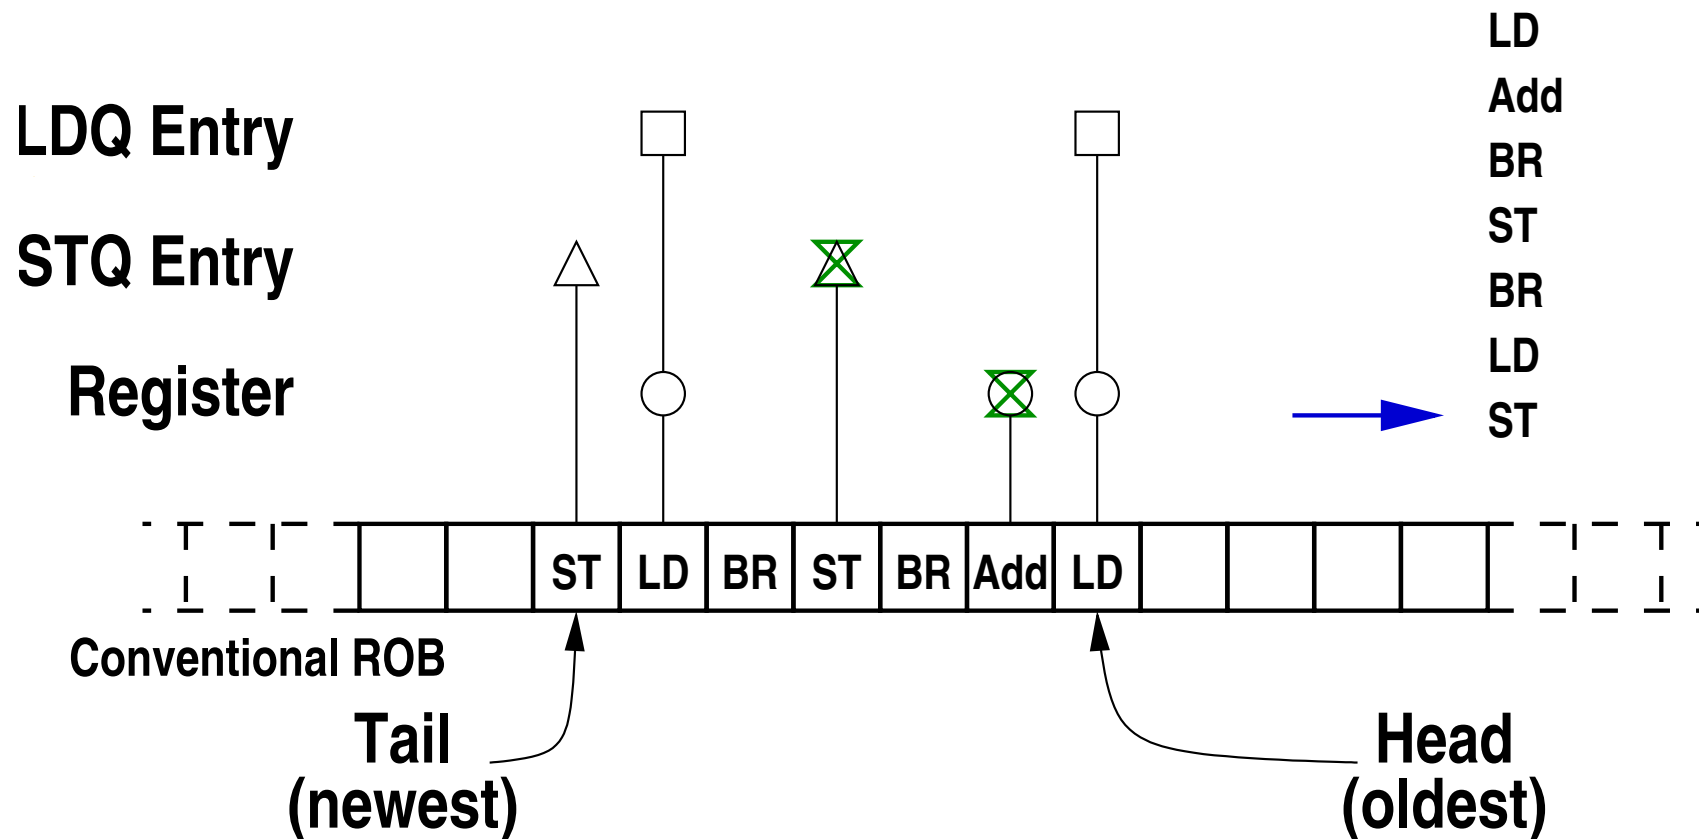
- Decouple resource recycling from retirement





EARLY RECYCLING

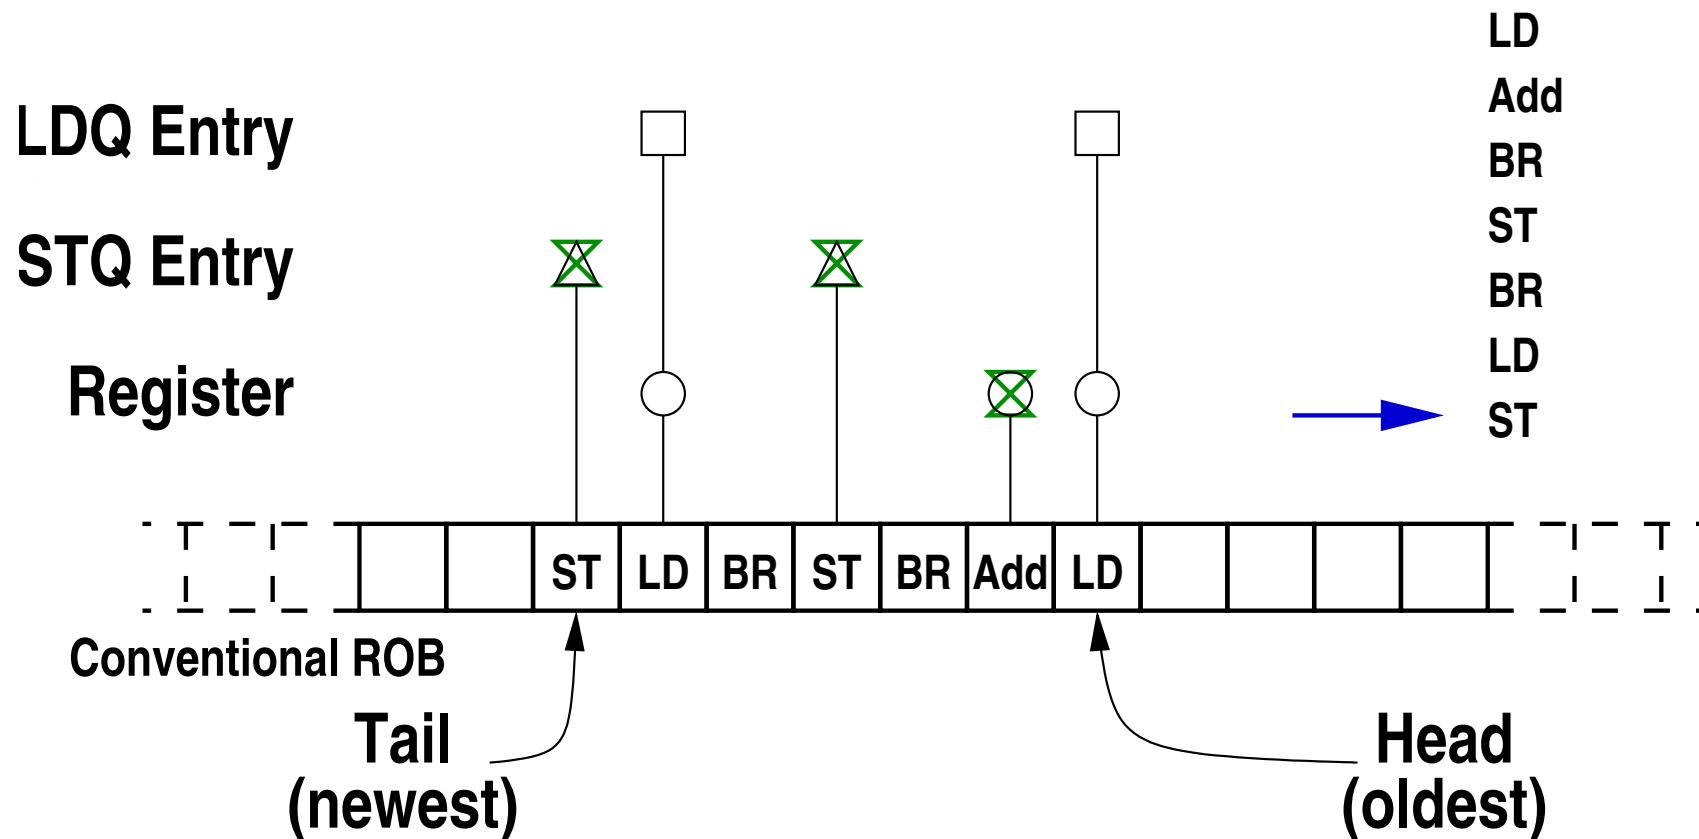
- Decouple resource recycling from retirement





EARLY RECYCLING

- Decouple resource recycling from retirement






CONSTRAINT: PRECISE STATE

- Events that require precise state:
 - Branch mispredictions
 - Memory replay traps
 - Exceptions
 - Interrupts
- Early recycling makes this difficult
 - E.g. recycled registers no longer available
 - E.g. memory may be overwritten prematurely





PROPOSAL: POINT OF NO RETURN

- Classify events according to frequency:
 - Frequent: branch mispredictions, memory replays
 - Infrequent: exceptions
- Split ROB into two logical sets of instructions
 - **Irreversible set**
 - Subject to infrequent events only
 - Checkpoint recovery
 - **Reversible set:**
 - Subject to frequent and infrequent events
 - Conventional ROB recovery mechanisms
 - Boundary: **Point of No Return (PNR)**

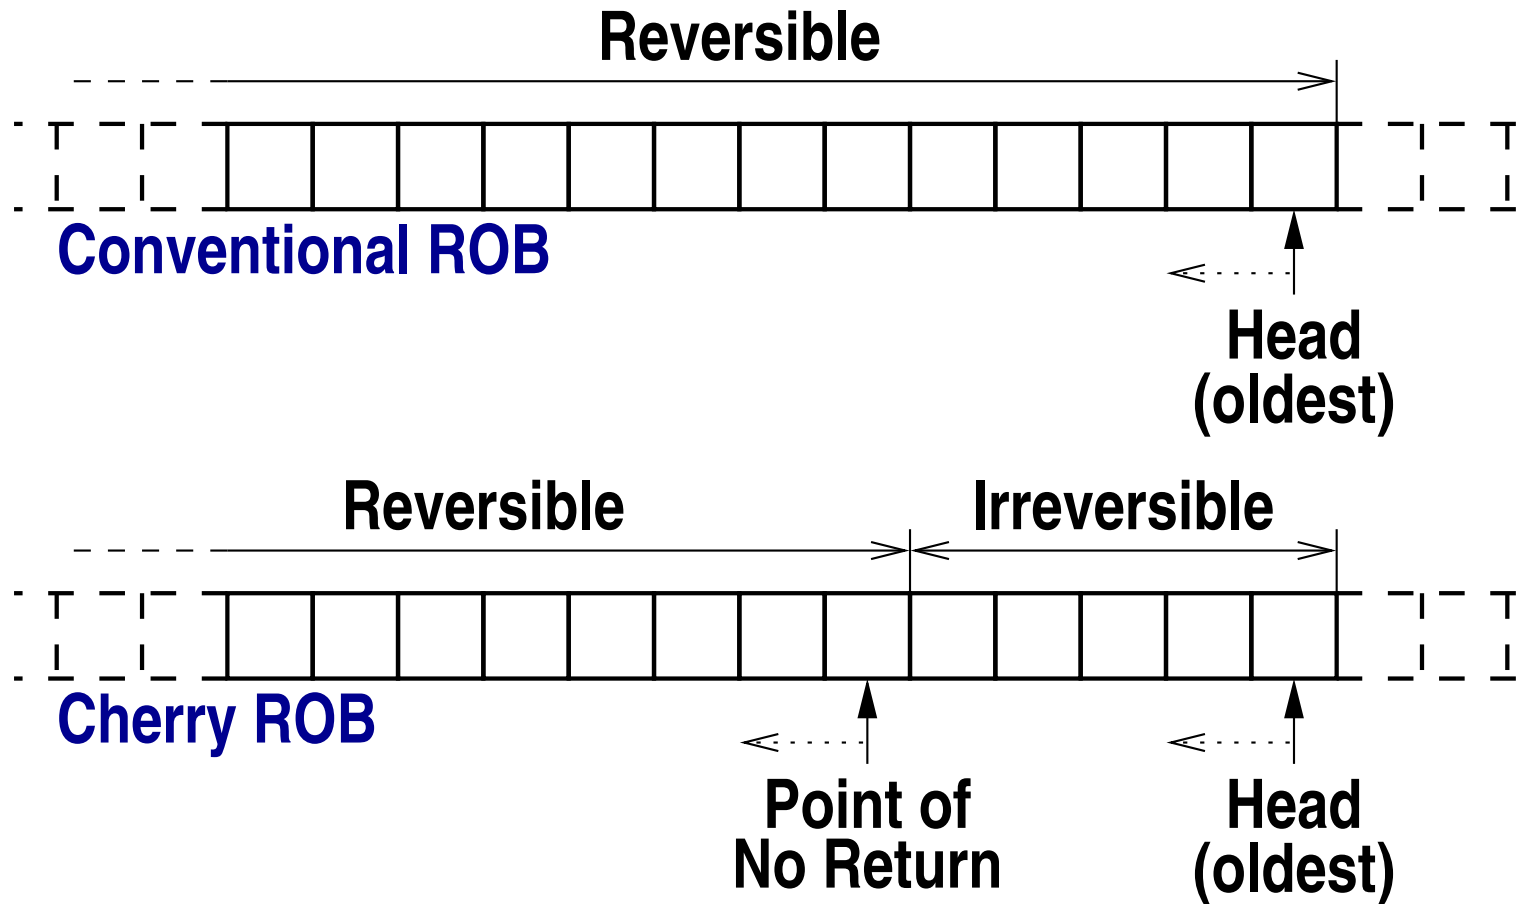



CONVENTIONAL VS CHERRY ROB



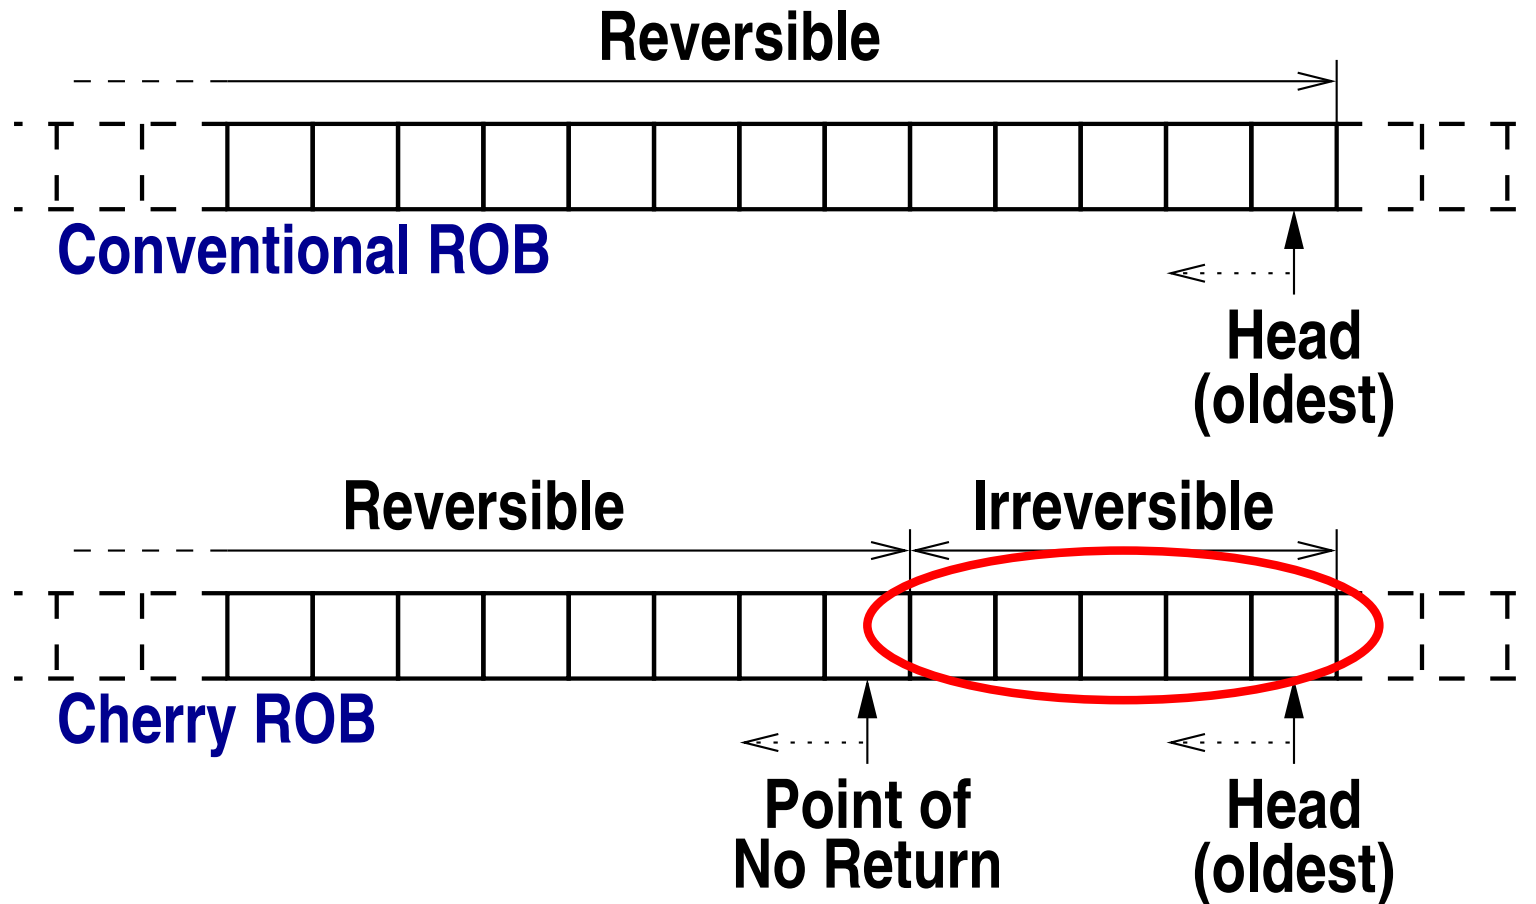


CONVENTIONAL VS CHERRY ROB






CONVENTIONAL VS CHERRY ROB





OUTLINE

- Motivation
- Overview
- Implementation
 - Checkpointing
 - Exception and interrupt handling
 - Early resource recycling
- Results




TIMELINE



Timeline




Conventional ROB




TIMELINE







TIMELINE






TIMELINE



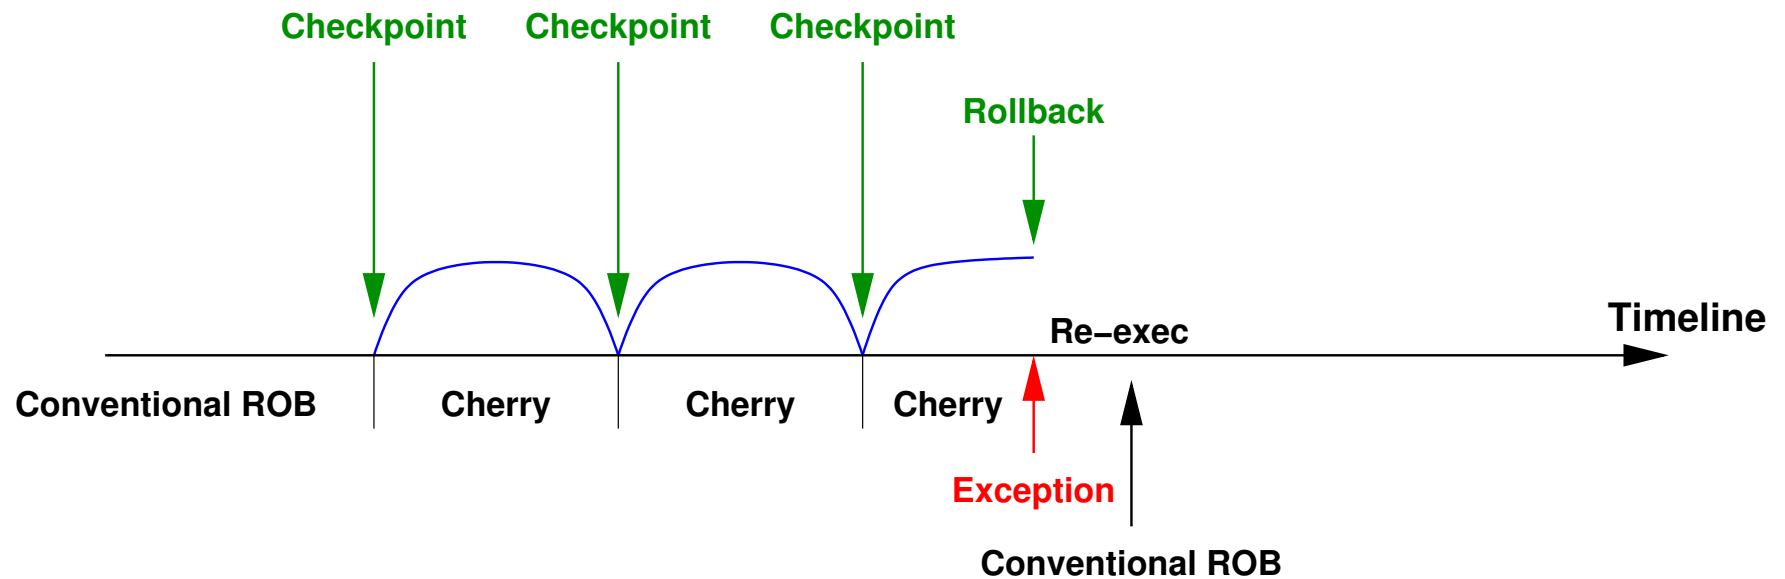



TIMELINE



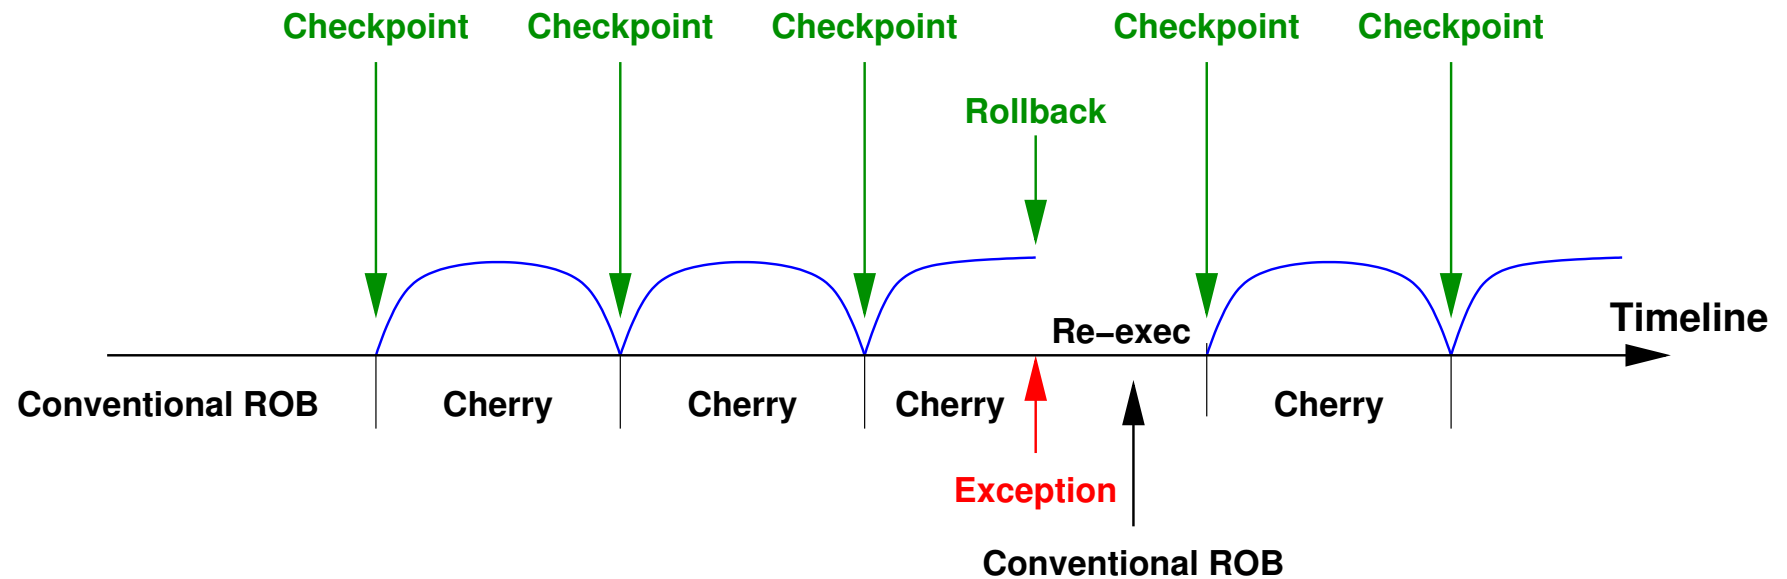



TIMELINE





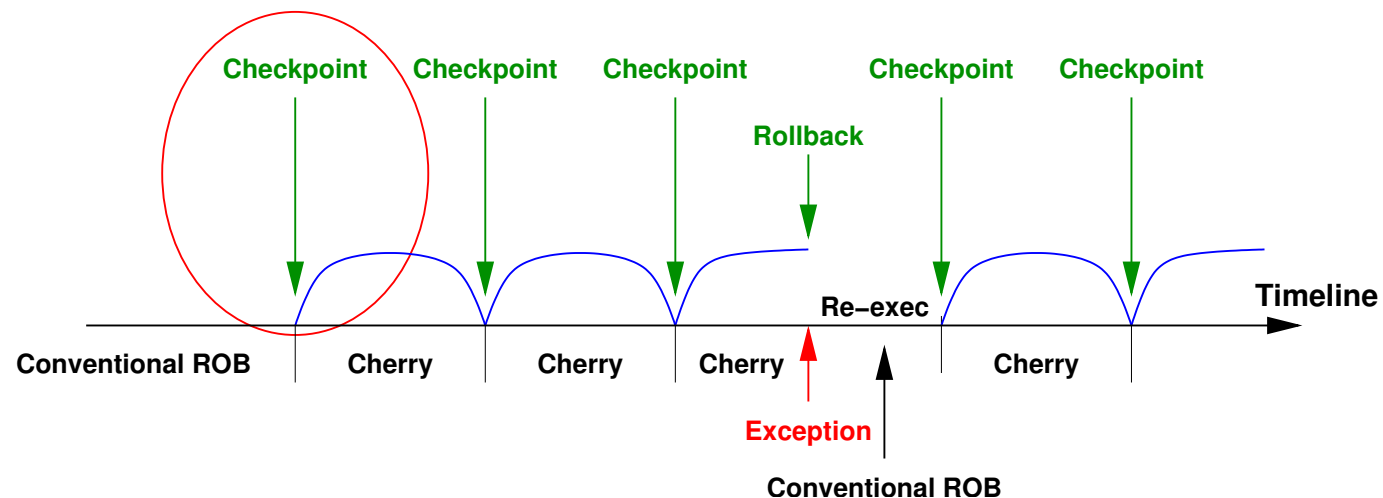
TIMELINE






ENTERING CHERRY MODE

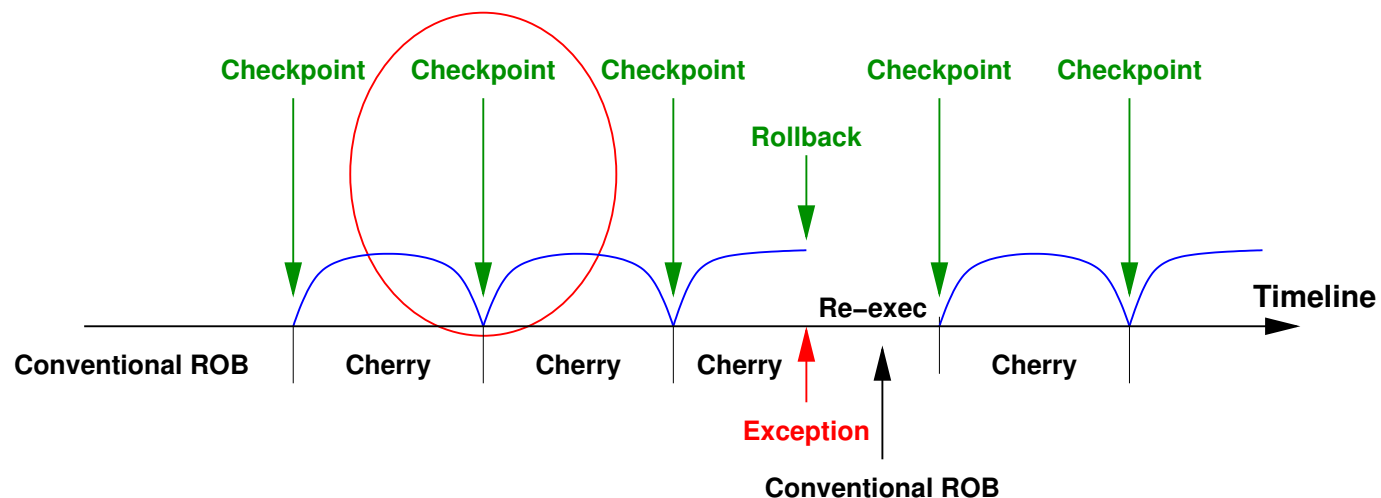
- Backup architectural registers
 - Can be done in a handful of cycles
- Allow PNR to race ahead of ROB head
- Updates to cache hierarchy set Volatile bit






PERIODIC CHECKPOINTING


- Done regularly to bound re-execution overhead
- Stop early recycling to create checkpoint (collapse)
 - Freeze PNR (stop recycling)
- Once ROB head catches up with PNR:
 - Clear Volatile bits from cache
 - Backup architectural registers






PRECISE EXCEPTION HANDLING


- Exception in Irreversible set:
 - Roll back to checkpoint
 - Temporarily disable recycling (for some time)
 - Re-execute in conventional OOO mode (w/o Cherry)
 - Allow exception to re-occur





PRECISE EXCEPTION HANDLING


- Exception in Irreversible set:
 - Roll back to checkpoint
 - Temporarily disable recycling (for some time)
 - Re-execute in conventional OOO mode (w/o Cherry)
 - Allow exception to re-occur






PRECISE EXCEPTION HANDLING


- Exception in Irreversible set:
 - Roll back to checkpoint
 - Temporarily disable recycling (for some time)
 - Re-execute in conventional OOO mode (w/o Cherry)
 - Allow exception to re-occur






PRECISE EXCEPTION HANDLING

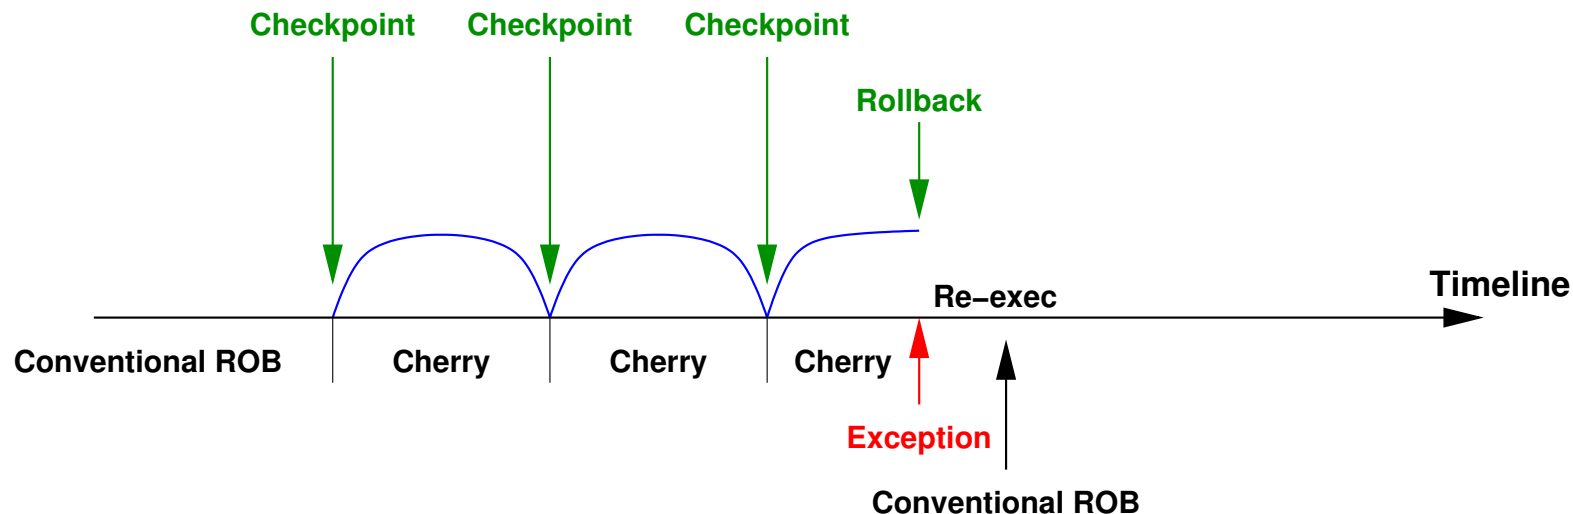
- Exception in Irreversible set:
 - Roll back to checkpoint
 - Temporarily disable recycling (for some time)
 - Re-execute in conventional OOO mode (w/o Cherry)
 - Allow exception to re-occur






PRECISE EXCEPTION HANDLING

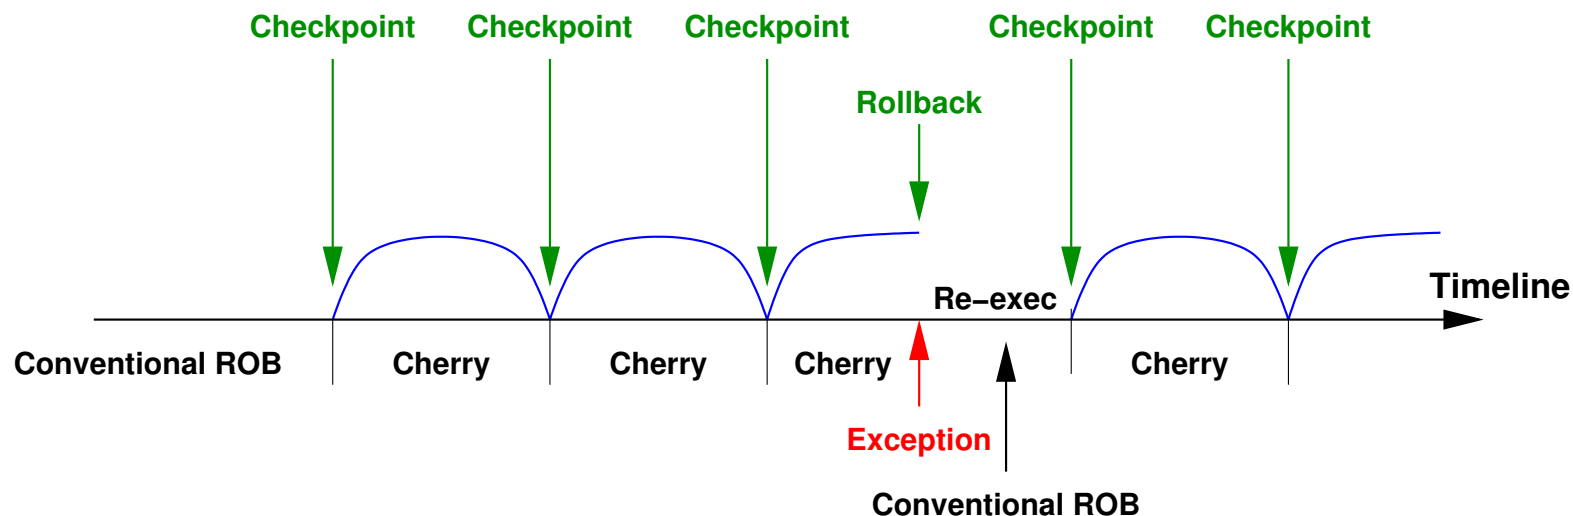
- Exception in Irreversible set:
 - Roll back to checkpoint
 - Temporarily disable recycling (for some time)
 - Re-execute in conventional OOO mode (w/o Cherry)
 - Allow exception to re-occur






PRECISE EXCEPTION HANDLING


- Exception in Irreversible set:
 - Roll back to checkpoint
 - Temporarily disable recycling (for some time)
 - Re-execute in conventional OOO mode (w/o Cherry)
 - Allow exception to re-occur






PRECISE EXCEPTION HANDLING


- Exception in Reversible set:
 - Disable resource recycling (freeze PNR)
 - Process exception
 - Take new checkpoint; re-enter Cherry mode






PRECISE EXCEPTION HANDLING


- Exception in Reversible set:
 - Disable resource recycling (freeze PNR)
 - Process exception
 - Take new checkpoint; re-enter Cherry mode






PRECISE EXCEPTION HANDLING


- Exception in Reversible set:
 - Disable resource recycling (freeze PNR)
 - Process exception
 - Take new checkpoint; re-enter Cherry mode





PRECISE EXCEPTION HANDLING


- Exception in Reversible set:
 - Disable resource recycling (freeze PNR)
 - Process exception
 - Take new checkpoint; re-enter Cherry mode






INTERRUPT HANDLING


- Interrupts are asynchronous: delay handling
 - Disable resource recycling (freeze PNR)
 - Process interrupt
 - Take new checkpoint and re-enter Cherry mode






INTERRUPT HANDLING


- Interrupts are asynchronous: delay handling
 - Disable resource recycling (freeze PNR)
 - Process interrupt
 - Take new checkpoint and re-enter Cherry mode






INTERRUPT HANDLING

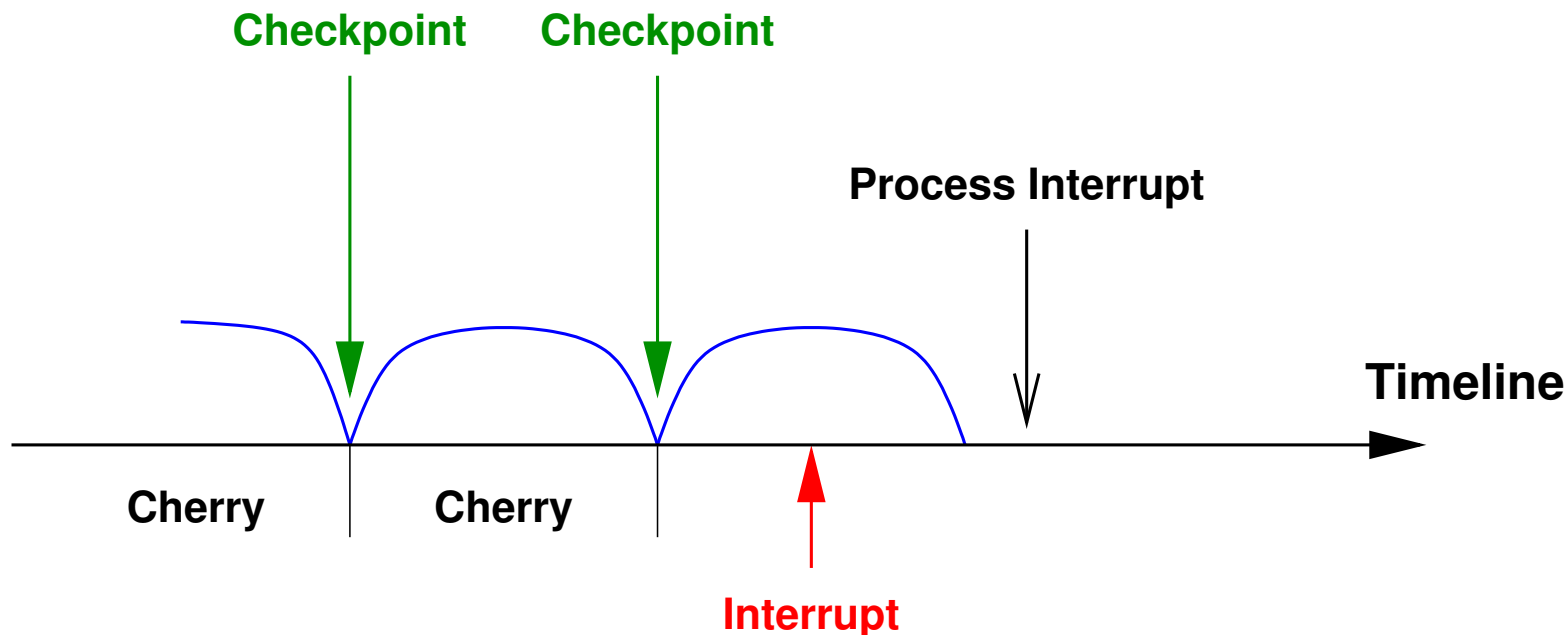
- Interrupts are asynchronous: delay handling
 - Disable resource recycling (freeze PNR)
 - Process interrupt
 - Take new checkpoint and re-enter Cherry mode






INTERRUPT HANDLING

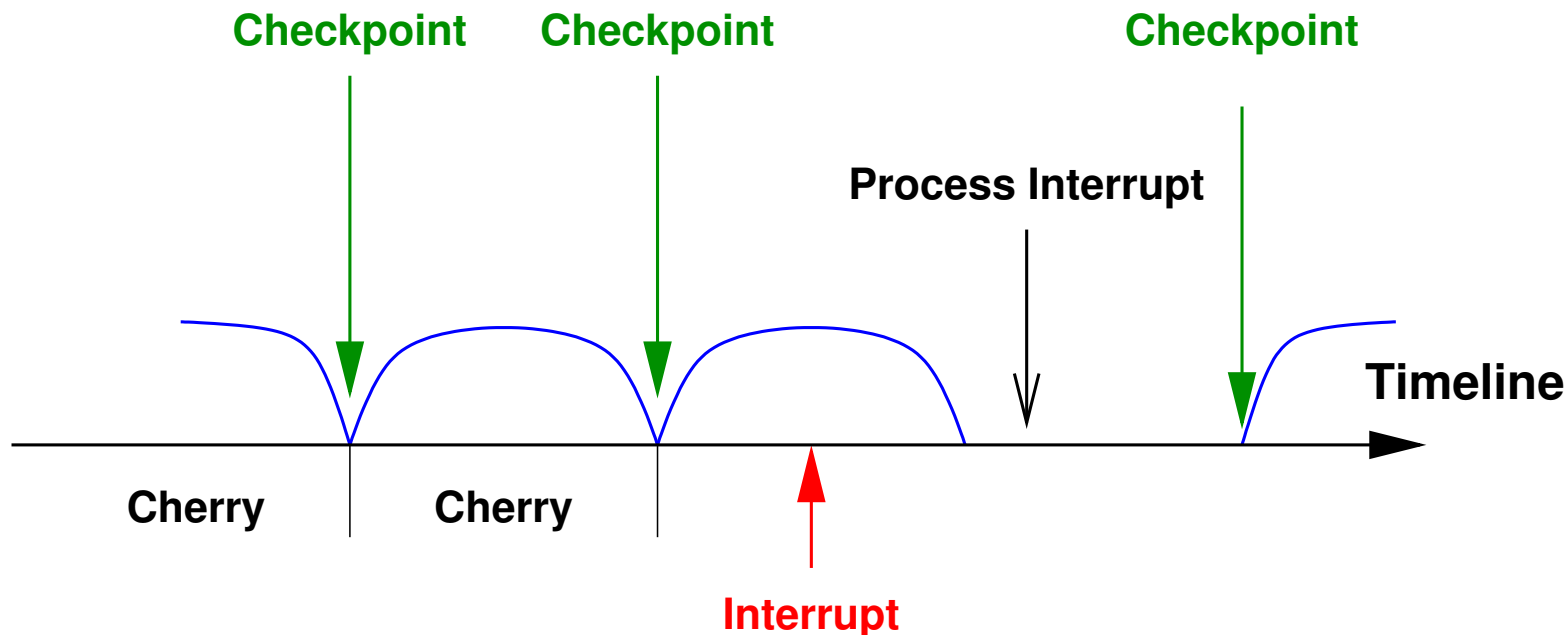
- Interrupts are asynchronous: delay handling
 - Disable resource recycling (freeze PNR)
 - Process interrupt
 - Take new checkpoint and re-enter Cherry mode





INTERRUPT HANDLING


- Interrupts are asynchronous: delay handling
 - Disable resource recycling (freeze PNR)
 - Process interrupt
 - Take new checkpoint and re-enter Cherry mode






OUTLINE


- Motivation
- Overview
- Implementation
 - Checkpointing
 - Exception and interrupt handling
 - Early resource recycling
- Results




DEFINING THE PNRs

- Each resource defines its own PNR
- Based on three ROB pointers:
 - U_b : oldest unresolved branch
 - U_s : oldest store with unresolved address
 - U_l : oldest load with unresolved address



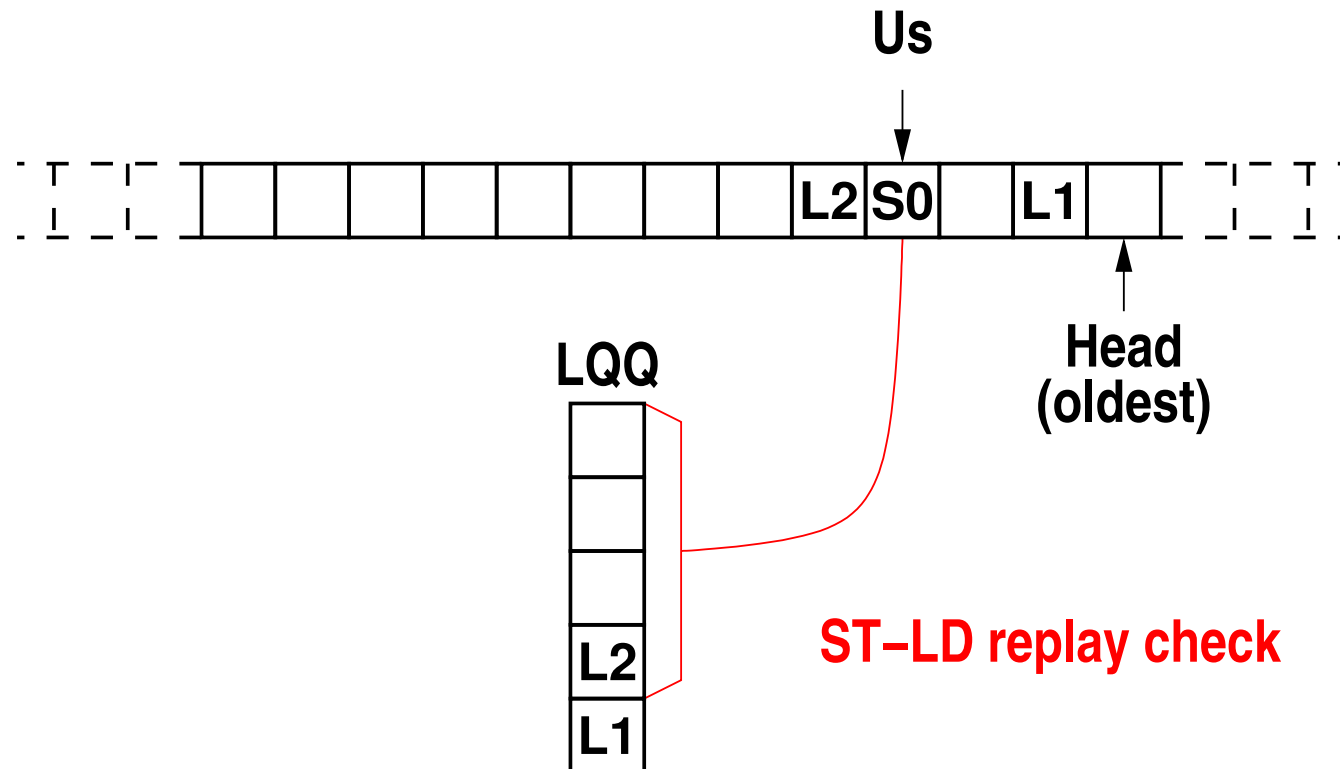



EARLY LDQ ENTRY RECYCLING





EARLY LDQ ENTRY RECYCLING






EARLY LDQ ENTRY RECYCLING

- Constraint: Must detect memory replay traps
 - $\text{older}(U_s)$: ST-LD replay trap
 - Not affected by branches

$$\text{PNR}_{\text{LDQ}} = U_s$$




EARLY REGISTER RECYCLING

- Constraint: Must not contain potentially useful data
 - $\text{older}(U_s)$: Free of ST-LD replay trap
 - $\text{older}(U_b)$: Free of branch misprediction

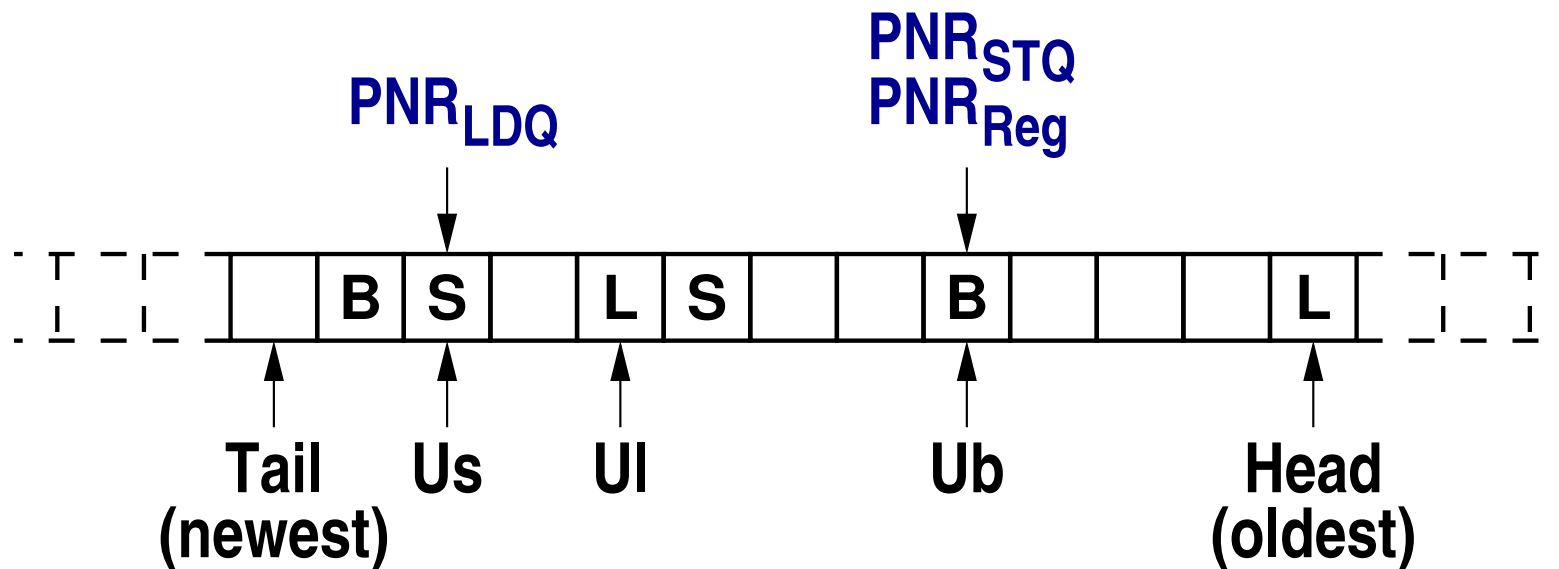
$$\text{PNR}_{\text{Reg}} = \text{oldest}(U_S, U_B)$$


- Of course register must be dead (paper)



EARLY RECYCLING SUMMARY


Resource	PNR Value
LDQ entries	U_S
STQ entries	$\text{oldest}(U_L, U_S, U_B)$
Registers	$\text{oldest}(U_S, U_B)$





EVALUATION

- Execution-driven simulation
- SPEC 2000 applications
- SGI MIPSPro compiler
- Compare Cherry against:
 - Processor with unlimited resources
 - Three aggressive base processor designs (paper)



PROCESSOR MODEL

Frequency: 3.2GHz

Fetch/issue/commit width: 8/8/12

I. window/ROB size: 128/384

Int/FP registers : 192/128

Ld/St units: 2/2

Int/FP/branch units: 7/5/3

Ld/St queue entries: 32/32

MSHRs: 24

Cherry checkpoint frequency: $5\mu s$

Up to 1 taken branch/cycle

RAS: 32 entries

BTB: 4K entries, 4-way assoc.


Branch predictor:

Hybrid with speculative update


Bimodal size: 8K entries


Two-level size: 64K entries

Cache	L1	L2	Bus & Memory
Size:	32KB	512KB	FSB frequency: 400MHz
RT:	2 cycles	10 cycles	FSB width: 128bit
Assoc:	4-way	8-way	Memory: 4-channel Rambus
Line size:	64B	128B	DRAM bandwidth: 6.4GB/s
Ports:	4	1	Memory RT: 120ns




OVERALL RESULTS






OVERALL RESULTS






OVERALL RESULTS






PNR ADVANCE


- PNR advance as a % of ROB use






PNR ADVANCE


- PNR advance as a % of ROB use






PNR ADVANCE


- PNR advance as a % of ROB use






PNR ADVANCE


- PNR advance as a % of ROB use





PNR ADVANCE


● PNR advance as a % of ROB use





SUMMARY

- Decoupling resource recycling from retirement
 - Split ROB into two logical sets
 - Use of ROB+checkpointing simultaneously
- Early recycling mechanism: LDQ, STQ, Registers
- Integration with speculative multithreading (paper)
- Speedup over baseline
 - 1.06 for SPECint2000
 - 1.26 for SPECfp2000



QUESTIONS?

José F. Martínez¹

Jose Renau²

Michael C. Huang³

Milos Prvulovic²

Josep Torrellas²



1



2




3





BACKUP SLIDES



RECYCLING PNRs

Resource	PNR Value
LDQ entries (uniprocessor)	U_S
LDQ entries (multiprocessor)	$\text{oldest}(U_L, U_S)$
STQ entries	$\text{oldest}(U_L, U_S, U_B)$
Registers (uniprocessor)	$\text{oldest}(U_S, U_B)$
Registers (multiprocessor)	$\text{oldest}(U_L, U_S, U_B)$



EARLY STQ ENTRY RECYCLING

- Constraint: Must not overwrite potentially useful data
 - $\text{older}(U_l)$: Unresolved LD may need old data
 - $\text{older}(U_s)$: No older LD is replayable
 - $\text{older}(U_b)$: Free of branch misprediction

$$\text{PNR}_{\text{STQ}} = \text{oldest}(U_L, U_S, U_B)$$



CHECKPOINT FREQUENCY


$$T_o = c_k + p_e c_e \quad (1)$$

$$p_e = \frac{T_c}{T_e} \quad (2)$$


$$c_e = s \frac{T_c}{2} + (s - 1) \quad (3)$$


$$\frac{T_o}{T_c} = \frac{c_k}{T_c} + \left(s - \frac{1}{2} \right) \frac{T_c}{T_e} \quad (4)$$

$$T_c = \sqrt{\frac{c_k T_e}{s - \frac{1}{2}}} \quad (5)$$




CHECKPOINT SELECTION






SPECINT2000 PARAMETERS

	Apps	Used ROB (%)	Irreversible Set (% of Used ROB)			Collapse Step (Cycles)
			Reg	LQ	SQ	
SPECint	bzip2	29.9	24.3	55.8	19.5	292.3
	crafty	28.8	33.4	97.6	28.6	41.9
	gcc	19.1	19.0	82.3	17.8	66.9
	gzip	28.5	65.5	81.7	8.5	47.1
	mcf	30.1	14.6	37.7	13.8	695.6
	parser	30.7	26.1	80.7	21.8	109.2
	perlbnk	12.2	24.6	89.9	20.5	23.3
	vortex	39.3	26.3	87.1	24.9	64.4
	vpr	32.9	25.2	83.6	21.5	165.1
	Average	27.9	28.7	77.4	19.7	167.3




SPECfp2000 PARAMETERS

	Apps	Used ROB (%)	Irreversible Set (% of Used ROB)			Collapse Step (Cycles)
			Reg	LQ	SQ	
SPECfp	applu	62.2	61.6	62.4	60.7	411.5
	apsi	76.8	82.3	83.1	81.6	921.1
	art	88.0	54.3	62.6	29.2	1247.3
	equake	41.6	61.6	69.1	57.3	135.2
	mesa	29.8	35.1	44.6	34.6	33.7
	mgrid	65.1	91.5	93.5	91.3	335.9
	swim	59.4	64.8	65.4	64.7	949.1
	wupwise	71.9	90.3	71.2	87.9	190.7
	Average	61.9	67.7	78.3	63.4	528.1




REGISTER ACCESS TIME






OTHER CONSIDERATIONS (PAPER)

- Exception on instruction older than youngest PNR:
 - Roll back to checkpoint
- Each PNR recycles at its own pace
- Not all PNRs critical, e.g. PNR_{LDQ}
- Early register recycling:
 - Must count in-flight consumers (Moudgill et al 1993)
- Optimal interval between checkpoints





CHERRY PARAMETERS

- Cherry checkpoint frequency: $5\mu s$
- Checkpoint overhead: 60 cycles ($18.75ns$)
- Avg time between exceptions: $800\mu s$ (not modeled)





INGORING U_b







PERFECT BRANCH PREDICTION






INTERRUPT COLLAPSING






INTERRUPT COLLAPSING


Interrupt






INTERRUPT COLLAPSING

Interrupt



Point of
No Return
& Head



EQUIVALENT SYSTEM

Resource	Baseline	SPECint2000	SPECfp2000
LDQ entries	32	80	128
STQ entries	32	64	112
Int Registers	192	240	240
FP Registers	128	128	192